

Economic Forecasting

Shrinkage methods

Sebastian Fossati

University of Alberta | E493 | 2023

- 1 Model selection
- 2 Ridge regression
- 3 The lasso
- 4 Ridge regression vs. lasso
- 5 Example: Used cars
- 6 A very large data set

Suppose we observe a dependent variable y and p potential explanatory variables (or predictors) x_1, x_2, \dots, x_p such that

$$y = f(x_1, x_2, \dots, x_p) + \varepsilon$$

Two main reasons to perform variable selection:

- *model interpretability*: including irrelevant variables leads to unnecessary complexity
- *prediction accuracy*: including irrelevant variables leads to less accurate predictions, specially when N is not much larger than p

The methods described before work well with small p ...

- but the number of models grows as p increases, with $p = 10$ there are $2^{10} = 1024$ possible models
- if $p < N$ but close to N , OLS will not be very accurate
- if $p > N$, OLS cannot be used!

Shrinkage methods

Shrinkage methods involve fitting a model with all p predictors (even if $p > N$).

- estimated coefficients are shrunk towards 0, some may be estimated to be exactly 0
- shrinkage reduces the variance of the estimators

- 1 Model selection
- 2 Ridge regression
- 3 The lasso
- 4 Ridge regression vs. lasso
- 5 Example: Used cars
- 6 A very large data set

Ridge regression

The **ridge regression estimates** $(\hat{\beta}^R)$ are given by

$$\hat{\beta}^R = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^N \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$

Remarks:

- that is, minimize $\text{SSR} + \lambda \sum_{j=1}^p \beta_j^2$
- $\lambda \geq 0$ is a *tuning parameter*
- $\lambda \sum_{j=1}^p \beta_j^2$ is the *shrinkage penalty*

Remarks:

- if $\lambda = 0$, $\hat{\beta}^R$ is the OLS estimate
- if $\lambda \rightarrow \infty$, $\hat{\beta}^R = \mathbf{0}$ (the null model)
- λ balances fit vs. shrinking of the estimates towards 0
- selecting a good value of λ is critical, more later

The ridge regression estimates can change substantially when a predictor is multiplied by a constant. Why?

As a result, it is better to first *standardize the predictors* such that they all have the same variance (= 1):

$$\bar{x}_{ij} = \frac{x_{ij}}{\sqrt{\frac{1}{N} \sum_{i=1}^N (x_{ij} - \bar{x}_j)^2}}.$$

Ridge regression vs. OLS

In general, if the relationship between y and the predictors is close to linear, OLS will have low bias but *may* have high variance.

Remarks:

- if p is close to N , OLS estimates will have very large variance
- if $p > N$, OLS cannot be used (variance is ∞)!
- ridge regression trades off an increase in bias for a decrease in variance (the effect of shrinkage)
- ridge regression works best when OLS estimates have high variance

Simulated example 1

True model: $y_i = \beta_1 x_{i1} + \beta_2 x_{i2} + \varepsilon_i, \quad i = 1, \dots, N$

And consider a small data set with 3 predictor variables such that:

- $\beta_1 = .75, \beta_2 = .50, \beta_j = 0$ for $j = 0, 3$
- $\sigma = 2$
- $N = 200$

Note that y is a function of x_1 and x_2 , all other variables are irrelevant.

- OLS is perfectly capable of handling this one
- but we may still get some insights on how ridge regression works

Simulated example 1

First, simulate data according to data generating process (DGP). Next, use `scale(x)` to standardize the variables.

```
# simulate data
n.obs <- 200
n.var <- 3
#
e <- rnorm(n.obs, mean = 0, sd = 2) # errors
x <- matrix(0, n.obs, n.var)
for (i in 1:n.var){ x[,i] <- rnorm(n.obs) }
#
y <- .75*x[,1] + .50*x[,2] + e
x <- scale(x)
data.all <- data.frame(x,y)
```

Simulated example 1

```
model1 <- lm(y ~ X1 + X2 + X3, data = data.all)
coeftest(model1)
```

```
##
## t test of coefficients:
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -0.100      0.143   -0.70   0.4853
## X1             0.983      0.144    6.83  1.1e-10 ***
## X2             0.387      0.144    2.69   0.0077 **
## X3             0.233      0.144    1.62   0.1069
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Simulated example 1

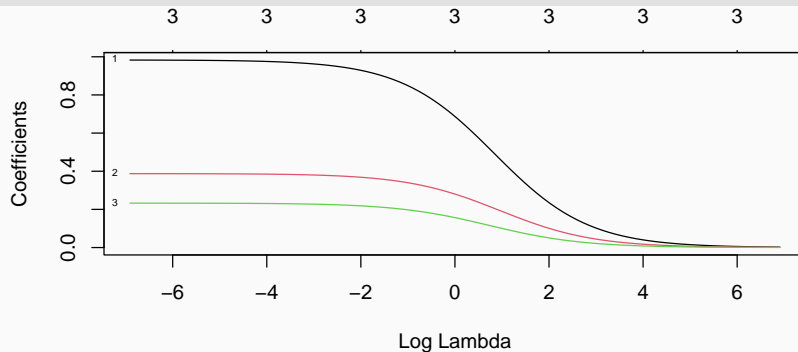
We will use the R package `glmnet`.

- `alpha=0` selects the ridge regression penalty, $\lambda \sum_{j=1}^p \beta_j^2$
- we are estimating the ridge regression for a range of values of λ ,
 $\lambda \in (.001, 1000)$

```
# ridge reg for different values of lambda  
grid <- 10^seq(3, -3, length = 100)  
ridge.mod <- glmnet(x, y, alpha = 0, lambda = grid)
```

Simulated example 1

```
# plot ridge results  
plot(ridge.mod, xvar = "lambda", label = TRUE)
```



- each line corresponds to a coefficient estimate as a function of λ (if $\lambda \rightarrow \infty$?)

Simulated example 1

We can compare the OLS and ridge estimates for a given λ .

- a large λ produces substantial shrinkage (relative to OLS)

```
# ridge
```

```
ridge.mod$lambda[1]
```

```
## [1] 1000
```

```
coef(ridge.mod)[,1]
```

```
## (Intercept)          V1          V2          V3
## -0.1002473    0.0022821    0.0010010    0.0004738
```

```
# least squares
```

```
model1$coef
```

```
## (Intercept)          X1          X2          X3
## -0.1002      0.9832      0.3875      0.2329
```


Simulated example 1

We can compare the OLS and ridge estimates for a given λ .

- a small λ produces less shrinkage

```
# ridge
```

```
ridge.mod$lambda[80]
```

```
## [1] 0.0163
```

```
coef(ridge.mod)[,80]
```

```
## (Intercept)      V1      V2      V3
##    -0.1002    0.9764    0.3851    0.2311
```

```
# least squares
```

```
model1$coef
```

```
## (Intercept)      X1      X2      X3
##    -0.1002    0.9832    0.3875    0.2329
```

- 1 Model selection
- 2 Ridge regression
- 3 The lasso
- 4 Ridge regression vs. lasso
- 5 Example: Used cars
- 6 A very large data set

The lasso

The **lasso estimates** ($\hat{\beta}^L$) are given by

$$\hat{\beta}^L = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^N \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j|$$

Remarks:

- that is, minimize $\text{SSR} + \lambda \sum_{j=1}^p |\beta_j|$
- $\lambda \geq 0$ is a *tuning parameter*
- $\lambda \sum_{j=1}^p |\beta_j|$ is the *shrinkage penalty*

Remarks:

- if $\lambda = 0$, $\hat{\beta}^L$ is the OLS estimate
- if $\lambda \rightarrow \infty$, $\hat{\beta}^L = \mathbf{0}$
- λ balances fit vs. shrinking of the estimates towards 0
- selecting a good value of λ is critical, more later

The lasso estimates can also change substantially when a predictor is multiplied by a constant. As a result, it is better to first standardize the predictors such that they all have the same variance (= 1).

In general, the lasso works very much like ridge regression.

But the lasso performs **variable selection**, ridge regression does not.

- ridge regression will include all p predictors in the final model
- the lasso may force some estimates to be exactly equal to 0
- if λ is sufficiently large, the ridge regression and lasso models may differ
- lasso models are generally much easier to interpret

Simulated example 1


We will use the R package `glmnet`.

- `alpha=1` selects the lasso regression penalty, $\lambda \sum_{j=1}^p |\beta_j|$
- we are estimating the lasso regression for a range of values of λ ,
 $\lambda \in (.001, 1000)$

```
# lasso reg for different values of lambda  
grid <- 10^seq(3, -3, length = 100)  
lasso.mod <- glmnet(x, y, alpha = 1, lambda = grid)
```

Simulated example 1

```
# plot ridge results  
plot(lasso.mod, xvar = "lambda", label = TRUE)
```



- each line corresponds to a coefficient estimate as a function of λ (if $\lambda \rightarrow \infty$?)

Simulated example 1

We can compare the OLS and lasso estimates for a given λ .

- a large λ produces substantial shrinkage (relative to OLS)

```
# ridge
```

```
lasso.mod$lambda[1]
```

```
## [1] 1000
```

```
coef(lasso.mod)[,1]
```

```
## (Intercept)          V1          V2          V3
##    -0.1002      0.0000      0.0000      0.0000
```

```
# least squares
```

```
model1$coef
```

```
## (Intercept)          X1          X2          X3
##    -0.1002      0.9832      0.3875      0.2329
```


Simulated example 1

We can compare the OLS and lasso estimates for a given λ .

- a small λ produces less shrinkage

```
# ridge
```

```
lasso.mod$lambda[80]
```

```
## [1] 0.0163
```

```
coef(lasso.mod)[,80]
```

```
## (Intercept)          V1          V2          V3
##    -0.1002      0.9674      0.3717      0.2160
```

```
# least squares
```

```
model1$coef
```

```
## (Intercept)          X1          X2          X3
##    -0.1002      0.9832      0.3875      0.2329
```

- 1 Model selection
- 2 Ridge regression
- 3 The lasso
- 4 Ridge regression vs. lasso
- 5 Example: Used cars
- 6 A very large data set

The ridge regression and lasso coefficients solve the problems:

$$\min_{\beta} \left\{ \sum_{i=1}^N \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \text{ subject to } \sum_{j=1}^p \beta_j^2 \leq s$$
$$\min_{\beta} \left\{ \sum_{i=1}^N \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \text{ subject to } \sum_{j=1}^p |\beta_j| \leq s$$

Remarks:

- OLS is the unconstrained solution
- lasso and ridge estimates are constrained solutions
- this reduces the variance as it keeps the estimates close to 0
- the “shape” of the restriction matters

Ridge regression vs. lasso

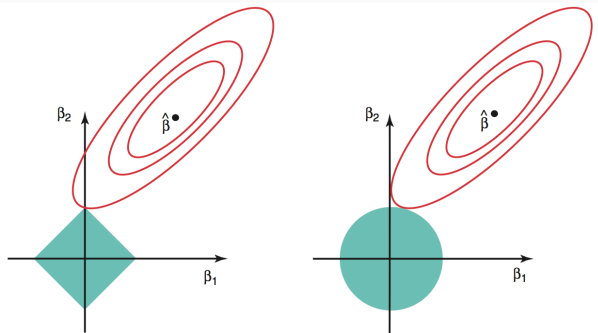


FIGURE 6.7. Contours of the error and constraint functions for the lasso (left) and ridge regression (right). The solid blue areas are the constraint regions, $|\beta_1| + |\beta_2| \leq s$ and $\beta_1^2 + \beta_2^2 \leq s$, while the red ellipses are the contours of the RSS.

- $\hat{\beta}$ is the (unrestricted) OLS solution, the lasso and ridge regression solutions may differ

Which is better?

- the lasso should perform better when a relatively small number of predictors are important and the remaining predictors have coefficients that are very small or equal to 0
- ridge regression should perform better when many predictors are (roughly) equally important
- cross-validation can be used to determine which approach is better in a particular data set

Simulated example 2

True model: $y_i = \beta_1 x_{i1} + \beta_2 x_{i2} + \varepsilon_i, \quad i = 1, \dots, N$

And consider a large data set of 50 mostly irrelevant predictor variables such that:

- $\beta_1 = .75, \beta_2 = .50, \beta_j = 0$ for $j = 0, 3, \dots, 50$
- $\sigma = 2$
- $N = 200$

Note that y is a function of x_1 and x_2 , all other variables are irrelevant.

- how can we find and estimate the model?

Simulated example 2

First, simulate data according to data generating process (DGP). Next, use `scale(x)` to standardize the variables.

```
# simulate data
n.obs <- 200
n.var <- 50
#
e <- rnorm(n.obs, mean = 0, sd = 2) # errors
x <- matrix(0, n.obs, n.var)
for (i in 1:n.var){ x[,i] <- rnorm(n.obs) }
#
y <- .75*x[,1] + .50*x[,2] + e
x <- scale(x)
data.all <- data.frame(x,y)
```



Simulated example 2

We are estimating the ridge and lasso regressions for a range of values of λ , $\lambda \in (.001, 1000)$.

```
#  
grid <- 10^seq(3, -3, length = 100)  
  
# ridge reg for different values of lambda  
ridge.mod <- glmnet(x, y, alpha = 0, lambda = grid)  
  
# lasso reg for different values of lambda  
lasso.mod <- glmnet(x, y, alpha = 1, lambda = grid)
```

Simulated example 2


```
# plot ridge results
plot(ridge.mod, xvar = "lambda", label = TRUE)
```



- each line corresponds to a coefficient estimate
- ridge does not perform variable selection

Simulated example 2

```
# plot lasso results  
plot(lasso.mod, xvar = "lambda", label = TRUE)
```



- each line corresponds to a coefficient estimate
- lasso performs variable selection

Selecting the tuning parameter

We can use cross-validation to select the tuning parameter λ .

Steps:

- split the sample into a *training set* and a *test set*
- use *k*-fold CV in the training set to find the value of λ that minimizes the cross-validation error
- compute the MSE for observations in the test set and evaluate the performance of the different models
- re-estimate using all available observations and the selected value of λ

Use the function `cv.glmnet` to find the λ value that minimizes the CV error.

Simulated example 2

```
# split the sample into train/test sets
train <- sample(nrow(data.all), round(nrow(data.all)/2))
cv.error <- rep(NA,3)

# least squares
ols.cv <- lm(y ~ x, data = data.all, subset = train)
cv.error[1] <- mean((y - predict(ols.cv, data.all))[-train]^2)


# ridge
ridge.cv <- cv.glmnet(x[train,], y[train], alpha = 0, nfolds = 10)
ridge.lam <- ridge.cv$lambda.min
cv.error[2] <- mean((y - predict(ridge.cv, s = ridge.lam, newx = x))[-train]^2)

# lasso
lasso.cv <- cv.glmnet(x[train,], y[train], alpha = 1, nfolds = 10)
lasso.lam <- lasso.cv$lambda.min
cv.error[3] <- mean((y - predict(lasso.cv, s = lasso.lam, newx = x))[-train]^2)
```

Simulated example 2

```
# plot ridge cv  
plot(ridge.cv)
```

50 50 50 50 50 50 50 50 50 50 50 50 50 50 50 50




- select the value of λ that minimizes the cross-validation error
- wide confidence intervals (estimate not very accurate)

Simulated example 2

```
# plot lasso cv
```

```
plot(lasso.cv)
```



- select the value of λ that minimizes the cross-validation error
- narrow confidence intervals (estimate more accurate)

Simulated example 2

The MSE for observations in the test set is reported below.

```
cv.error <- data.frame(cv.error)
rownames(cv.error) <- c("ols", "ridge", "lasso")
cv.error
```

```
##      cv.error
## ols      8.015
## ridge    4.497
## lasso    4.758
```

Remarks:

- OLS performs poorly (relative to the other methods)
- ridge and lasso work best in this data set (why?)

Simulated example 2

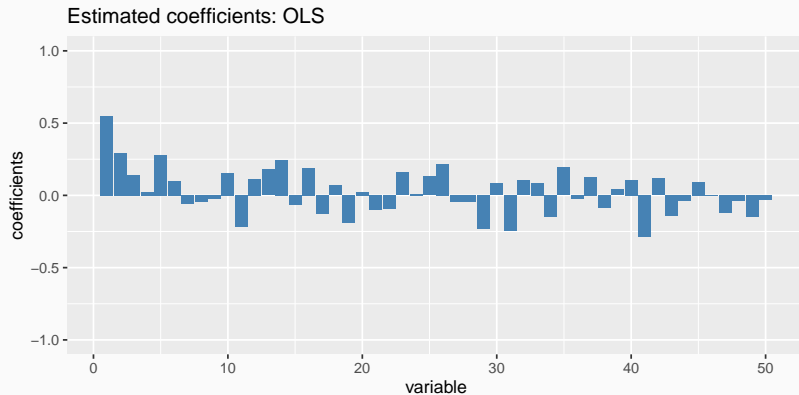
Re-estimate regression using all the available observations and the selected value of λ .

```
# least squares coefficients
model1 <- lm(y ~ x, data = data.all)
coef.ls <- data.frame("x" = seq(1:n.var), "beta" = coef(model1)[-1])

# ridge coefficients
model2 <- glmnet(x, y, alpha = 0, lambda = ridge.lam)
coef.ridge <- data.frame("x" = seq(1:n.var), "beta" = coef(model2)[-1])


# lasso coefficients
model3 <- glmnet(x, y, alpha = 1, lambda = lasso.lam)
coef.lasso <- data.frame("x" = seq(1:n.var), "beta" = coef(model3)[-1])
```

Simulated example 2



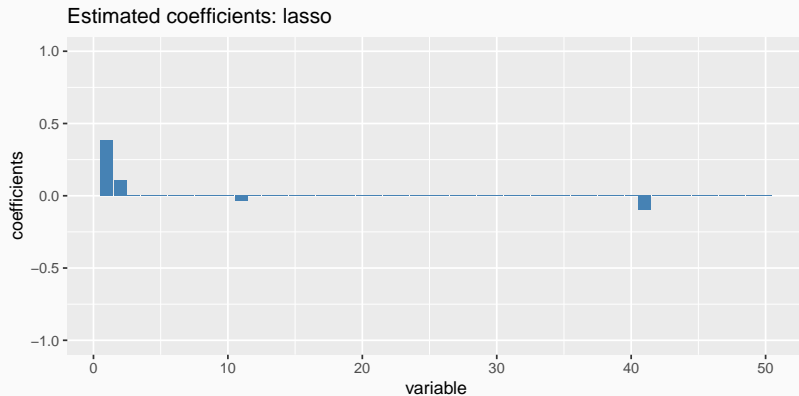
■ no shrinkage, all (most) coefficients different from 0

Simulated example 2



- noticeable shrinkage, all (most) coefficients different from 0

Simulated example 2



■ most coefficients shrink to 0

Since lasso performs variable selection, we can re-estimate the model using all available observations and only the selected variables.

Least squares estimates are less biased than lasso estimates.

Simulated example 2

```
# estimate the lasso selected model
```

```
coeftest(lm(y ~ X1 + X2 + X11 + X41, data = data.all))
```

```
##  
## t test of coefficients:  
##  
##           Estimate Std. Error t value Pr(>|t|)  
## (Intercept) -0.0292    0.1368   -0.21   0.831  
## X1           0.6134    0.1382    4.44 1.5e-05 ***  
## X2           0.3142    0.1381    2.27  0.024 *  
## X11          -0.2587    0.1382   -1.87  0.063 .  
## X41          -0.3447    0.1382   -2.49  0.013 *  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- 1 Model selection
- 2 Ridge regression
- 3 The lasso
- 4 Ridge regression vs. lasso
- 5 Example: Used cars
- 6 A very large data set

Example: Used cars

Suppose you want to sell your car of a certain make, type, year, miles, condition and other features.

Prediction analysis can help you uncover the average advertised price of cars with similar characteristics.

- that helps decide what price you may want to put on your ad

Example: Used cars

Consider a sample of offers for used Toyota Camry cars in 2018 in Chicago (Békés and Kézdi, 2021).

Data:

- source: scraped from a website
- characteristics: year of make (age), odometer (miles), etc.
- data cleaning: drop erroneous observations, hybrid cars, trucks...

Example: Used cars

Load data set and filter offers from the Chicago area.

```
# read data
data <-
  read.csv(
    "data/used_cars_work.csv",
    header = TRUE,
    stringsAsFactors = TRUE
  )
# focus only on Chicago
data <- data %>%
  filter(area == "chicago")
```

Example: Used cars

We will use the following predictors:


- age: measuring how old is the car (continuous, linear)
- odometer: measuring miles the car traveled (continuous, linear)
- car type: LE, XLE, SE (missing in about 30% of the observations, factor-set of dummies, incl. N/A)
- condition: good condition, excellent condition, or it is like new (missing for about one third of the ads, factor-set of dummies, incl. N/A)
- car's engine has 6 cylinders (20% of ads say this; 43% says 4 cylinders, and the rest has no information, binary for 6 cylinders)

Example: Used cars

Perform **best subset selection** using all the variables in Model 4 allowing for up to 10 variables.

```
# select variables
x <- data %>%
  dplyr::select(
    age, agesq, odometer, odometersq, SE, LE, XLE,
    cond_likenew, cond_excellent, cond_good, cylind6
  ) %>%
  data.matrix()
#
price <- data$price
data2 <- data.frame(x, price)
#
regfit.all <- regsubsets(price~. , data = data2, nvmax = 10)
reg.summary <- summary(regfit.all)
```

Example: Used cars

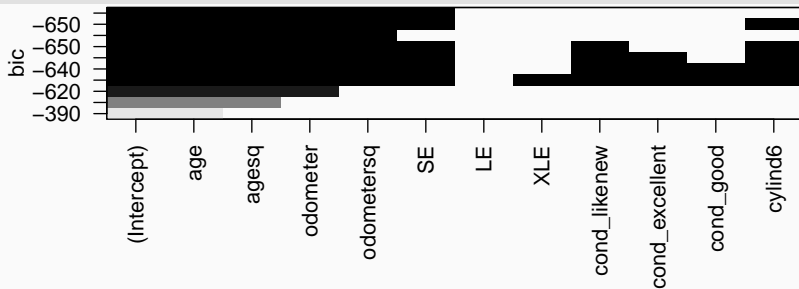


■ BIC and AIC (C_p) select $k = 5$

Example: Used cars

```
# plot variable
```

```
plot(regfit.all, scale = "bic")
```

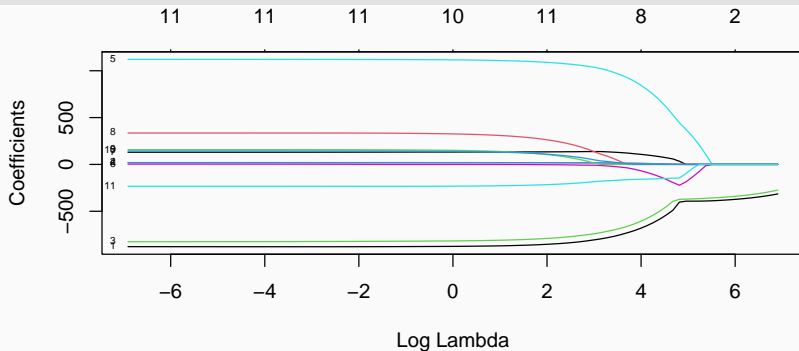


Example: Used cars

```
# lasso reg for different values of lambda
```

```
lasso.mod <- glmnet(x, price, alpha = 1, lambda = 10^seq(3, -3, length = 100))
```

```
plot(lasso.mod, xvar = "lambda", label = TRUE)
```



Example: Used cars

A more detailed look (prices in levels).

##	ols	bss	lasso
## (Intercept)	19948.916	20035.02	19764.700
## age	-876.959	-873.74	-855.876
## agesq	18.653	18.42	17.837
## odometer	-826.489	-832.01	-805.685
## odometersq	19.832	20.12	18.929
## SE	1122.004	1173.09	1101.982
## LE	3.256	NA	-2.465
## XLE	130.122	NA	134.158
## cond_likenew	334.838	NA	286.443
## cond_excellent	157.129	NA	121.354
## cond_good	143.040	NA	120.532
## cylind6	-233.506	NA	-224.179

Example: Used cars

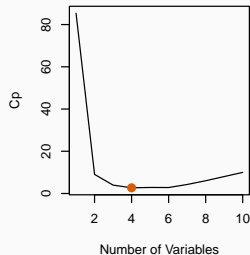
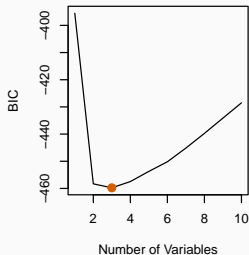
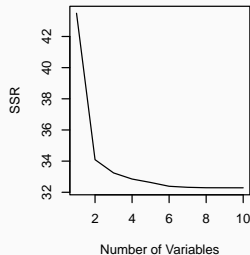
Should we work in levels or **logs**?

```
# select variables
x <- data %>%
  dplyr::select(
    age, agesq, odometer, odometersq, SE, LE, XLE,
    cond_likenew, cond_excellent, cond_good, cylind6
  ) %>%
  data.matrix()

#
lnprice <- data$lnprice
data2 <- data.frame(x, lnprice)

#
regfit.all <- regsubsets(lnprice~. , data = data2, nvmax = 10)
reg.summary <- summary(regfit.all)
```

Example: Used cars




■ BIC selects $k = 3$, AIC (Cp) selects $k = 4$

Example: Used cars

```
# plot variable
```

```
plot(regfit.all, scale = "bic")
```

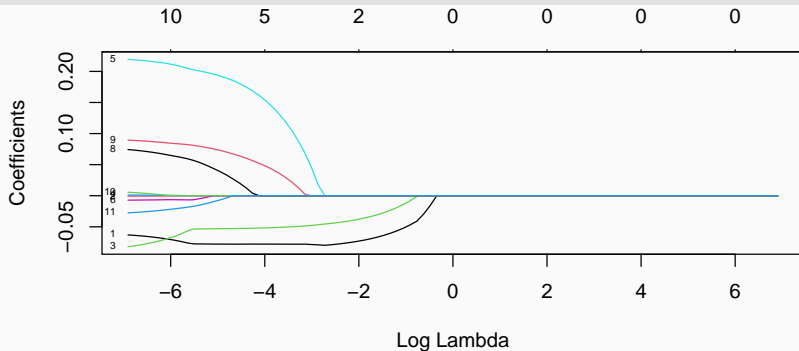


Example: Used cars

```
# lasso reg for different values of lambda
```

```
lasso.mod <- glmnet(x, lnprice, alpha = 1, lambda = 10^seq(3, -3, length = 100))
```

```
plot(lasso.mod, xvar = "lambda", label = TRUE)
```



Example: Used cars

A more detailed look (prices in logs).

##	ols	bss	lasso
## (Intercept)	10.1355346	10.10577	10.0318223
## age	-0.0563904	-0.08429	-0.0778367
## agesq	-0.0009484	NA	-0.0001926
## odometer	-0.0935311	-0.05372	-0.0521044
## odometersq	0.0015117	NA	0.0000000
## SE	0.2247556	0.22733	0.1634097
## LE	-0.0080826	NA	0.0000000
## XLE	-0.0057798	NA	0.0000000
## cond_likenew	0.0808348	NA	0.0000000
## cond_excellent	0.0927394	NA	0.0532273
## cond_good	0.0090029	NA	0.0000000
## cylind6	-0.0311883	NA	0.0000000

- 1 Model selection
- 2 Ridge regression
- 3 The lasso
- 4 Ridge regression vs. lasso
- 5 Example: Used cars
- 6 A very large data set**

Simulated example 3

True model: $y_i = \beta_1 x_{i1} + \beta_2 x_{i2} + \varepsilon_i, \quad i = 1, \dots, N$

And consider a large data set of 500 mostly irrelevant predictor variables such that:

- $\beta_1 = .75, \beta_2 = .50, \beta_j = 0$ for $j = 3, \dots, 500$
- $\sigma = 2$
- $N = 200$

Note that y is a function of x_1 and x_2 , all other variables are irrelevant.

- how can we find and estimate the model?

Simulated example 3

First, simulate data according to data generating process (DGP). Next, use `scale(x)` to standardize the variables.


```
# simulate data
n.obs <- 200
n.var <- 500
#
e <- rnorm(n.obs, mean = 0, sd = 2) # errors
x <- matrix(0, n.obs, n.var)
for (i in 1:n.var){ x[,i] <- rnorm(n.obs) }
#
y <- .75*x[,1] + .50*x[,2] + e
x <- scale(x)
data.all <- data.frame(x,y)
```


Simulated example 3

```
# lasso reg for different values of lambda
```


```
lasso.mod <- glmnet(x, y, alpha = 1, lambda = 10seq(3, -3, length = 100))
```

```
plot(lasso.mod, xvar = "lambda", label = TRUE)
```

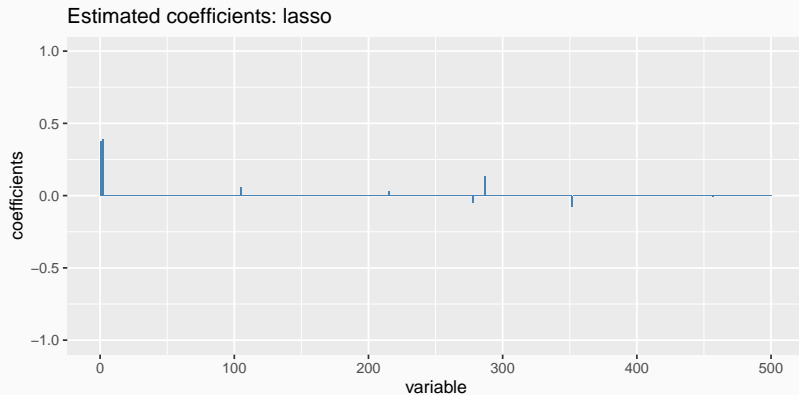


Simulated example 3

```
lasso.cv <- cv.glmnet(x, y, alpha = 1, nfolds = 10)  
plot(lasso.cv)
```



Simulated example 3



Simulated example 3


What would we find if we run 500 single OLS regressions?

Simulated example 3

What would we find if we run 500 single OLS regressions?

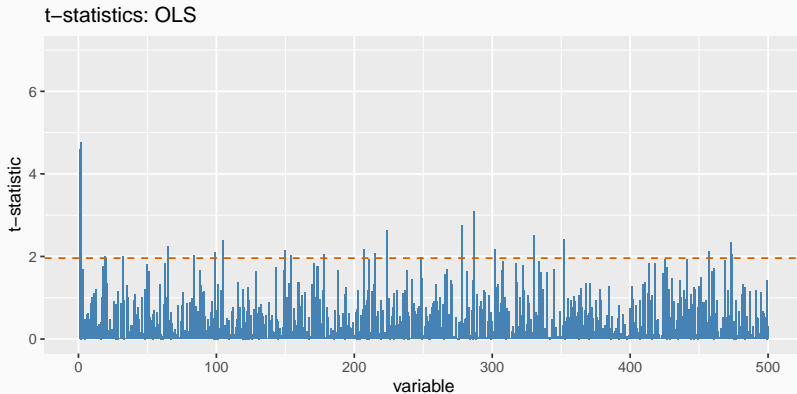
```
#
coef.ols <- rep(NA, n.var)
test.ols <- rep(NA, n.var)
for (i in 1:n.var){
  tmp0 <- lm(y ~ x[,i])
  coef.ols[i] <- coeftest(tmp0)[2,1] # estimated slope
  test.ols[i] <- abs(coeftest(tmp0)[2,3]) # t-stat
}
# single regression coefficients
coef.ols <-
  data.frame(
    "x" = seq(1:n.var),
    "beta" = coef.ols,
    "test" = test.ols
  )
```

Simulated example 3



■ x_1 and x_2 show stronger effect but many coefficients $\neq 0$

Simulated example 3



■ many t-statistics > 2 !