

# Economic Forecasting

Forecasting with time series data


---

Sebastian Fossati


University of Alberta | E493 | 2023

- 1 Some simple forecasting methods
- 2 Transformations and adjustments
- 3 Residual diagnostics
- 4 Evaluating forecast accuracy
- 5 Prediction intervals

# How would you forecast this time series?



# How would you forecast this time series?



## Some simple forecasting methods

Some simple forecasting methods are surprisingly effective.

Here are four methods that we will use as benchmarks for other forecasting methods.

- 1 Average method
- 2 Naïve method
- 3 Seasonal naïve method
- 4 Drift method

## Average method

- forecast of all future values is equal to mean of historical data  $\{y_1, \dots, y_T\}$
- forecasts:  $\hat{y}_{T+h|T} = \bar{y} = (y_1 + \dots + y_T)/T$

```
# y contains the time series  
# h is the forecast horizon  
meanf(y, h)
```

## Naïve method

- forecasts equal to last observed value
- consequence of efficient market hypothesis
- forecasts:  $\hat{y}_{T+h|T} = y_T$

```
# y contains the time series  
# h is the forecast horizon  
naive(y, h) # alternative rwf(y, h)
```

### Seasonal naïve method

- forecasts equal to last value from same season (eg, the same month of the previous year)
- forecasts:  $\hat{y}_{T+h|T} = y_{T+h-m(k+1)}$ , where  $m$  = seasonal period and  $k$  is the integer part of  $(h - 1)/m$

```
# y contains the time series  
# h is the forecast horizon  
snaive(y, h)
```



## Drift method

- forecasts equal to last value plus average change
- equivalent to extrapolating a line drawn between first and last observations
- forecasts:


$$\hat{y}_{T+h|T} = y_T + \frac{h}{T-1} \sum_{t=2}^T (y_t - y_{t-1})$$

```
# y contains the time series  
# h is the forecast horizon  
rwf(y, h, drift = TRUE)
```

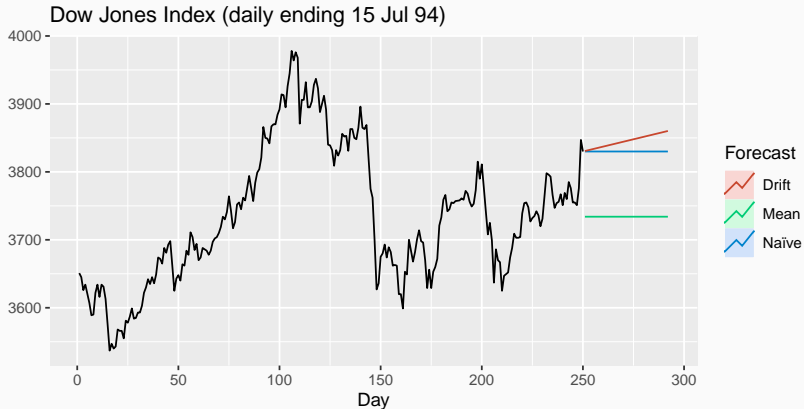
## Example 1: quarterly beer production

```
# get data
beer2 <- window(ausbeer, start = 1992, end = c(2007,4))
# get forecasts
fcst.m <- meanf(beer2, h = 10)
fcst.n <- naive(beer2, h = 10)
fcst.s <- snaive(beer2, h = 10)
# plot
autoplot(beer2) +
  autolayer(fcst.m, PI = FALSE, series = "Mean") +
  autolayer(fcst.n, PI = FALSE, series = "Naïve") +
  autolayer(fcst.s, PI = FALSE, series = "Seasonal naïve") +
  ggtitle("Forecasts for quarterly beer production") +
  xlab("Year") + ylab("Megalitres") +
  guides(colour = guide_legend(title = "Forecast"))
```

## Example 1: quarterly beer production



## Example 2: Dow Jones Index



- 1 Some simple forecasting methods
- 2 Transformations and adjustments
- 3 Residual diagnostics
- 4 Evaluating forecast accuracy
- 5 Prediction intervals

## Transformations and adjustments

Adjusting the data can simplify the patterns in the historical data, making the pattern more consistent across the whole sample.

Examples:

- 1 mathematical transformations
- 2 calendar adjustments
- 3 population adjustments
- 4 inflation adjustments

Simpler patterns usually lead to more accurate forecasts!

If the data show different variation at different levels of the series, then a transformation can be useful.

Logarithms, in particular, are useful because they are more interpretable: changes in a log value are **relative (percent) changes on the original scale**.

Denote original observations as  $y_1, \dots, y_T$  and transformed observations as  $w_1, \dots, w_T$ .


## Mathematical transformations

Square root:  $w_t = \sqrt{y_t}$


Logarithm:  $w_t = \log(y_t)$



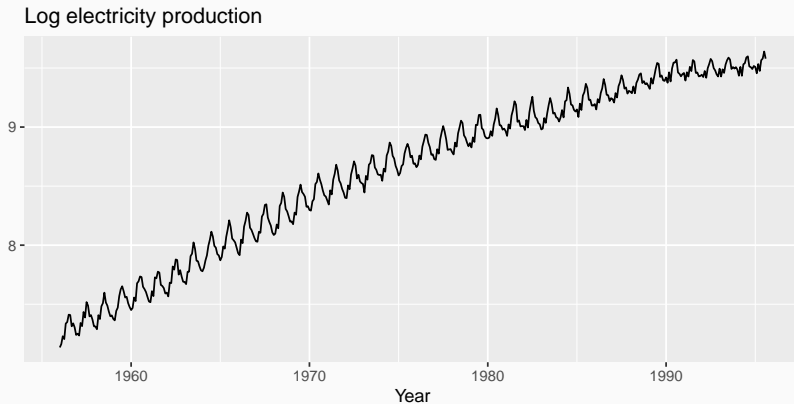
# Variance stabilization



# Variance stabilization



# Variance stabilization



## Box-Cox transformations

Each of these transformations is close to a member of the family of **Box-Cox transformations**:


$$w_t = \begin{cases} \log(y_t), & \lambda = 0 \\ (y_t^\lambda - 1)/\lambda, & \lambda \neq 0 \end{cases}$$

- $\lambda = 1$ : (no substantive transformation)
- $\lambda = \frac{1}{2}$ : (square root plus linear transformation)
- $\lambda = 0$ : (natural logarithm)

# Box-Cox transformations

```
# box-cox
```

```
autoplot(BoxCox(elec, lambda = 1/3))
```



Remarks:

- $y_t^\lambda$  for  $\lambda$  close to zero behaves like logs
- if some  $y_t \leq 0$ , no power transformation is possible unless all  $y_t$  adjusted by **adding a constant to all values**
- simple values of  $\lambda$  are easier to explain
- often no transformation ( $\lambda = 1$ ) needed
- transformation can have very large effect on PI
- choosing  $\lambda = 0$  is a simple way to force forecasts to be positive

## Automated Box-Cox transformations

```
# box-cox selection
```

```
BoxCox.lambda(elec)
```

```
## [1] 0.2654
```

- this attempts to balance the seasonal fluctuations and random variation across the series
- a low value of  $\lambda$  can give extremely large prediction intervals

We must reverse the transformation (or *back-transform*) to obtain forecasts on the original scale. The reverse Box-Cox transformations are given by


$$y_t = \begin{cases} \exp(w_t), & \lambda = 0 \\ (\lambda W_t + 1)^{1/\lambda}, & \lambda \neq 0 \end{cases}$$



# Back-transformation

```
# forecasting  
fit <- snaive(elec, lambda = 1/3)  
autoplot(fit)
```

Forecasts from Seasonal naive method




# Back-transformation

```
# a closer look
```

```
autoplot(fit, include = 120)
```

Forecasts from Seasonal naive method



Other adjustments may include:

- calendar adjustments: some variation seen in seasonal data may be due to simple calendar effects (eg, different numbers of days in each month)
- population adjustments: data affected by population changes can be adjusted to give per-capita data
- inflation adjustments: data affected by the value of money are best adjusted before modelling

- 1 Some simple forecasting methods
- 2 Transformations and adjustments
- 3 Residual diagnostics**
- 4 Evaluating forecast accuracy
- 5 Prediction intervals

The fitted value  $\hat{y}_{t|t-1}$  is the forecast of  $y_t$  based on observations  $y_1, \dots, y_{t-1}$ .

- sometimes drop the subscript:  $\hat{y}_t \equiv \hat{y}_{t|t-1}$
- not true forecasts if parameters are estimated on all data

For example:

- $\hat{y}_t = \bar{y}$  for average method

# Forecasting residuals

Difference between observed value and its fitted value:

$$e_t = y_t - \hat{y}_{t|t-1}.$$

## Assumptions

- 1  $\{e_t\}$  uncorrelated (if they aren't, information left in residuals should be used in computing forecasts)
- 2  $\{e_t\}$  have mean zero (if they don't, forecasts are biased)

# Forecasting residuals

Difference between observed value and its fitted value:

$$e_t = y_t - \hat{y}_{t|t-1}.$$


## Assumptions

- 1  $\{e_t\}$  uncorrelated (if they aren't, information left in residuals should be used in computing forecasts)
- 2  $\{e_t\}$  have mean zero (if they don't, forecasts are biased)

## Useful properties (for prediction intervals)

- 3  $\{e_t\}$  have constant variance
- 4  $\{e_t\}$  are normally distributed

## Example: Google stock price





Naïve forecast:

$$\hat{y}_{t|t-1} = y_{t-1}$$

Naïve forecast:

$$\hat{y}_{t|t-1} = y_{t-1}$$

$$e_t = y_t - y_{t-1}$$


Naïve forecast:

$$\hat{y}_{t|t-1} = y_{t-1}$$

$$e_t = y_t - y_{t-1}$$


Note:  $e_t$  are one-step-forecast residuals.

## Example: Google stock price



## Example: Google stock price

```
res <- residuals(naive(goog200))  
autoplot(res) + xlab("Day") + ylab("")
```




## Properties:

- we *expect* residuals to look like white noise (uncorrelated, mean zero, constant variance)
- so a standard residual diagnostic is to check the ACF of the residuals of a forecasting method
- we can also check other properties such as normality

## Example: Google stock price


```
# histogram
```

```
gghistogram(res, add.normal = TRUE) + ggtitle("Histogram of residuals")
```



## Example: Google stock price

```
# acf of residuals  
ggAcf(res) + ggtitle("ACF of residuals")
```





Consider a *whole set* of autocorrelations ( $r_k$  values), and develop a test to see whether the set is significantly different from a zero set.

### Box-Pierce test

$$Q = T \sum_{k=1}^h r_k^2$$

where  $h$  is max lag being considered and  $T$  is number of observations.

- if each  $r_k$  close to zero,  $Q$  will be **small**
- if some  $r_k$  away from zero,  $Q$  will be **large**
- about  $h = 10$  or  $h = 2m$  for seasonal data

### Ljung-Box test

$$Q^* = T(T+2) \sum_{k=1}^h (T-k)^{-1} r_k^2$$

where  $h$  is max lag being considered and  $T$  is number of observations.

- about  $h = 10$  or  $h = 2m$  for seasonal data
- better performance, especially in small samples

For the Google example:

```
# test
Box.test(res, lag = 10, fitdf = 0, type = "Lj")

##
##  Box-Ljung test
##
## data:  res
## X-squared = 11, df = 10, p-value = 0.4
```




- if data are WN,  $Q^*$  has  $\chi^2$  distribution with degrees of freedom  $(h - \# \text{ parameters in model})$
- when applied to raw data, set  $\# \text{ parameters} = 0$

## checkresiduals function

*# useful function!*

```
checkresiduals(naive(goog200), test = FALSE)
```

Residuals from Naive method



## checkresiduals function

```
# useful function!  
checkresiduals(naive(goog200), plot = FALSE)  
  
##  
##  Ljung-Box test  
##  
## data:  Residuals from Naive method  
## Q* = 11, df = 10, p-value = 0.4  
##  
## Model df: 0.    Total lags used: 10
```

- 1 Some simple forecasting methods
- 2 Transformations and adjustments
- 3 Residual diagnostics
- 4 Evaluating forecast accuracy**
- 5 Prediction intervals

A perfect fit can always be obtained by using a model with enough parameters (over-fitting).

- over-fitting a model to data is just as bad as failing to identify a systematic pattern in the data
- a common practice is to separate the available data into a training set and a test set



# Training and test sets



$$\text{MAE} = \text{mean}(|e_{T+h}|)$$

$$\text{MSE} = \text{mean}(e_{T+h}^2)$$

$$\text{RMSE} = \sqrt{\text{mean}(e_{T+h}^2)}$$

$$\text{MAPE} = 100\text{mean}(|e_{T+h}|/|y_{T+h}|)$$

# Training and test sets



$$\text{MAE} = \text{mean}(|e_{T+h}|)$$

$$\text{MSE} = \text{mean}(e_{T+h}^2)$$


$$\text{RMSE} = \sqrt{\text{mean}(e_{T+h}^2)}$$

$$\text{MAPE} = 100\text{mean}(|e_{T+h}|/|y_{T+h}|)$$

Remarks:

- MAE, MSE, RMSE are all scale dependent
- MAPE is scale independent but is only sensible if  $y_t \gg 0$  for all  $t$

## Example 1: quarterly beer production



## Example 1: quarterly beer production

```
beer2 <- window(ausbeer, start = 1992, end = c(2007,4)) # training set
beer3 <- window(ausbeer, start = 2008) # test set
beerfit1 <- meanf(beer2, h = 10)
beerfit2 <- rwf(beer2, h = 10)
beerfit3 <- snaive(beer2, h = 10)
accuracy(beerfit1, beer3)
accuracy(beerfit2, beer3)
accuracy(beerfit3, beer3)
```

	RMSE	MAE	MAPE
Mean method	38.45	34.83	8.28
Naïve method	62.69	57.40	14.18
Seasonal naïve method	14.31	13.40	3.17

To evaluate the **out-of-sample** performance of one or more models we can construct a series of  $h$ -step ahead predictions with the forecasting horizon  $h$  fixed.

For example, we can construct a series of  $R$  1-step ahead predictions following these steps:

- estimate the model with  $t = 1, \dots, T$  and construct  $\hat{Y}_{T+1|T}$
- re-estimate the model with  $t = 1, \dots, T + 1$  and construct  $\hat{Y}_{T+2|T+1}$
- repeat until  $t = 1, \dots, T + R - 1$  and  $\hat{Y}_{T+R|T+R-1}$

# Time series cross-validation

## Traditional evaluation




# Time series cross-validation

## Traditional evaluation



## Time series cross-validation




# Time series cross-validation

## Traditional evaluation



## Time series cross-validation



- forecast accuracy averaged over test sets
- an “evaluation on a rolling forecasting origin”
- or “pseudo out-of-sample forecast evaluation”



## Example 1: quarterly beer production

```
# data and set up
beer1 <- window(ausbeer, start = 1992, end = c(2009,4))
n.end <- 2004.75 # 2004Q4
# set matrix for storage, 20 obs in test set
pred <- matrix(rep(NA,80),20,4)
# loop
for(i in 1:20){
  tmp0 <- 1992
  tmp1 <- n.end+(i-1)*.25
  tmp <- window(beer1, tmp0, tmp1)
  pred[i,1] <- window(beer1, tmp1+.25, tmp1+.25) # actual
  # compute forecasts
  pred[i,2] <- meanf(tmp, h = 1)$mean
  pred[i,3] <- rwf(tmp, h = 1)$mean
  pred[i,4] <- snaive(tmp, h = 1)$mean
}
```

## Example 1: quarterly beer production

```
# compute rmse
beerfit1 <- pred[,1] - pred[,2]
beerfit2 <- pred[,1] - pred[,3]
beerfit3 <- pred[,1] - pred[,4]
rmse1 <- sqrt(mean(beerfit1^2, na.rm=TRUE))
rmse2 <- sqrt(mean(beerfit2^2, na.rm=TRUE))
rmse3 <- sqrt(mean(beerfit3^2, na.rm=TRUE))
# display rmse
cbind(rmse1,rmse2,rmse3)

##          rmse1 rmse2 rmse3
## [1,] 37.41 51.44 13.26
```

- choose forecasting model with the smallest RMSE computed using time series cross-validation

## tsCV function

```
# time series cross-validation function
beerfit1 <- tsCV(beer1, meanf, h = 1)
beerfit2 <- tsCV(beer1, rwf, h = 1)
beerfit3 <- tsCV(beer1, snaive, h = 1)

# compute rmse
rmse1 <- sqrt(mean(beerfit1[52:71]^2))
rmse2 <- sqrt(mean(beerfit2[52:71]^2))
rmse3 <- sqrt(mean(beerfit3[52:71]^2))


# display rmse
cbind(rmse1,rmse2,rmse3)

##          rmse1 rmse2 rmse3
## [1,] 37.41 51.44 13.26
```

## Example 1: quarterly beer production

```
# organize results
```

```
fcast.all <- ts(pred, start = 2005, freq = 4, names = c("Actual", "Mean", "Naive", "SNaive"))  
autoplot(fcast.all)
```



- 1 Some simple forecasting methods
- 2 Transformations and adjustments
- 3 Residual diagnostics
- 4 Evaluating forecast accuracy
- 5 Prediction intervals

## Prediction intervals

A forecast  $\hat{y}_{T+h|T}$  is (usually) the mean of the conditional distribution  $y_{T+h} \mid y_1, \dots, y_T$ .

- a prediction interval gives a region within which we expect  $y_{T+h}$  to lie with a specified probability

Assuming forecast errors are normally distributed, then a 95% PI is

$$\hat{y}_{T+h|T} \pm 1.96\hat{\sigma}_h$$

where  $\hat{\sigma}_h$  is the st dev of the  $h$ -step distribution.

### Remarks:

- point forecasts are often useless without prediction intervals
- prediction intervals require a stochastic model (with random errors, etc)
- multi-step forecasts for time series require a more sophisticated approach (with PI getting wider as the forecast horizon increases)

## Prediction intervals

Assume residuals are normal, uncorrelated,  $\text{sd} = \hat{\sigma}$ :

**Mean forecasts:**  $\hat{\sigma}_h = \hat{\sigma} \sqrt{1 + 1/T}$

**Naïve forecasts:**  $\hat{\sigma}_h = \hat{\sigma} \sqrt{h}$

**Seasonal naïve forecasts:**  $\hat{\sigma}_h = \hat{\sigma} \sqrt{k + 1}$

**Drift forecasts:**  $\hat{\sigma}_h = \hat{\sigma} \sqrt{h(1 + h/T)}$

where  $k$  is the integer part of  $(h - 1)/m$ .

Note that when  $h = 1$  and  $T$  is large, these all give the same approximate value  $\hat{\sigma}$ .



When  $h = 1$ ,  $\hat{\sigma}_h$  can be estimated from the residuals.

```
# compute sigma hat
res <- residuals(naive(goog200))
res_sd <- sqrt(mean(res^2, na.rm=TRUE))
# compute 05% interval
c(tail(goog200,1)) + 1.96 * res_sd * c(-1,1)

## [1] 519.3 543.6
```

## Prediction intervals

```
# predictions and prediction intervals  
naive(goog200, level = 95)
```


##	Point Forecast	Lo 95	Hi 95
## 201	531.5	519.3	543.6
## 202	531.5	514.3	548.7
## 203	531.5	510.4	552.6
## 204	531.5	507.1	555.8
## 205	531.5	504.3	558.7
## 206	531.5	501.7	561.3
## 207	531.5	499.3	563.7
## 208	531.5	497.1	565.9
## 209	531.5	495.0	568.0
## 210	531.5	493.0	570.0

# Prediction intervals

```
# plot prediction intervals
```

```
autoplot(naive(goog200))
```

Forecasts from Naive method



### Remarks:

- computed automatically using: `naive()`, `snaive()`, `rwf()`, `meanf()`, etc
- use `level` argument to control coverage
- check residual assumptions before believing them
- usually too narrow due to unaccounted uncertainty